

Dissolve^{struct}— A Library for Distributed Structured Prediction

Tribhuvanesh Orekondy

ETH Zurich

TOREKOND@STUDENT.ETHZ.CH

Martin Jaggi

ETH Zurich

JAGGI@INF.ETHZ.CH

Aurelien Lucchi

ETH Zurich

AURELIEN.LUCCHI@INF.ETHZ.CH

Editor: t.b.d.

Abstract

This paper describes DISSOLVE^{struct}, a modular and flexible open source software package for distributed training of structured prediction models, such as structured SVMs. Project website: github.com/dalab/dissolve.

We support a broad range of applications, and interfaces to scala, java and python. Our framework is empowered by the fault tolerant SPARK computing platform, and automatically adopts to the existing tradeoffs of computation vs communication cost on real world systems. The proposed distributed algorithm combines the recent communication efficient CoCoA scheme (Jaggi et al., 2014) with the state of the art primal-dual structured prediction solvers (Lacoste-Julien et al., 2013), and improves further by adding some new ideas for caching oracle answers. The framework allows approximate inference, and provides a similar standard interface as SVM^{struct} for the user.

Keywords: Structured Prediction, Structured SVM, Distributed Training

1. Introduction

Structured prediction has gained a lot of popularity over the past few years due to its ability to process structured objects such as images or text documents.

The structured support vector machine (SSVM) is a particularly successful variant of this approach that can be optimized in various ways, including a cutting-plane, stochastic gradient descent or a primal-dual scheme such as MARTIN’s paper. The latter is especially appropriate for large-scale problems ...

Remainder omitted in this sample. See <http://www.jmlr.org/papers/> for full paper.

2. Structured SVM formulation

A structured model predicts the labeling Y for a given input X by maximizing some score function $S_{\mathbf{w}} : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$, i.e.,

$$\hat{Y} = \operatorname{argmax}_{Y \in \mathcal{Y}} S_{\mathbf{w}}(Y) = \operatorname{argmax}_{Y \in \mathcal{Y}} \mathbf{w}^T \Psi(X, Y), \quad (1)$$

where \mathbf{w} are the model parameters and $\Psi(X, Y)$ is the *feature map* corresponding to the input X and the labeling Y .

Given a set of N training examples $\mathcal{D} = \{(X^1, Y^1), \dots, (X^N, Y^N)\}$ where $X^i \in \mathcal{X}$ and $Y^i \in \mathcal{Y}$ is the associated labeling, the learning task consists in finding model parameters \mathbf{w} that achieve low empirical loss subject to some regularization. In other words, we seek

$$\begin{aligned} \mathbf{w}^* &= \underset{\mathbf{w}}{\operatorname{argmin}} \mathcal{L}(\mathcal{D}, \mathbf{w}) \\ &= \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{n=1}^N l(X^n, Y^n, \mathbf{w}) + R(\mathbf{w}), \end{aligned} \quad (2)$$

where l is the *surrogate loss* function, a quantity that is usually related to and often upper-bounds the training error, and $R(\mathbf{w})$ is the regularizer (such as the L2 norm of \mathbf{w}) that helps prevent overfitting. The most common choice of l is the hinge loss, as used in Taskar et al. (2003); Tsochantaridis et al. (2005), which is defined as

$$l(Y^n, Y_{\mathbf{w}}^*, \mathbf{w}) = [S_{\mathbf{w}}(Y_{\mathbf{w}}^*) + \Delta(Y^n, Y_{\mathbf{w}}^*) - S_{\mathbf{w}}(Y^n)]_+ \quad (3)$$

The most popular method to solve problems of the form (2) is the stochastic subgradient method (SGD) TOOD: REF. Although this method has been quite successful it requires tuning parameters to achieve good performance (SAY SOMETHING about parallel SGD). We instead rely the CoCoA framework (Jaggi et al., 2014). WRITE DOWN DUAL and explain how step-size is selected...

The Lagrange dual of the above n -slack-formulation (??) has $m := \sum_i |\mathcal{Y}_i|$ variables or potential ‘support vectors’. Writing $\alpha_i(\mathbf{y})$ for the dual variable associated with the training example i and potential output $\mathbf{y} \in \mathcal{Y}_i$, the dual problem is given by

$$\begin{aligned} \min_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^m \\ \alpha \geq 0}} \quad & f(\boldsymbol{\alpha}) := \frac{\lambda}{2} \|A\boldsymbol{\alpha}\|^2 - \mathbf{b}^T \boldsymbol{\alpha} \\ \text{s.t.} \quad & \sum_{\mathbf{y} \in \mathcal{Y}_i} \alpha_i(\mathbf{y}) = 1 \quad \forall i \in [n], \end{aligned} \quad (4)$$

where the matrix $A \in \mathbb{R}^{d \times m}$ consists of the m columns $A := \left\{ \frac{1}{\lambda n} \boldsymbol{\psi}_i(\mathbf{y}) \in \mathbb{R}^d \mid i \in [n], \mathbf{y} \in \mathcal{Y}_i \right\}$, and the vector $\mathbf{b} \in \mathbb{R}^m$ is given by $\mathbf{b} := \left(\frac{1}{n} L_i(\mathbf{y}) \right)_{i \in [n], \mathbf{y} \in \mathcal{Y}_i}$. Given a dual variable vector $\boldsymbol{\alpha}$, we can use the Karush-Kuhn-Tucker optimality conditions to obtain the corresponding primal variables $\mathbf{w} = A\boldsymbol{\alpha} = \sum_{i, \mathbf{y} \in \mathcal{Y}_i} \alpha_i(\mathbf{y}) \frac{\boldsymbol{\psi}_i(\mathbf{y})}{\lambda n}$, see Appendix ?? . The gradient of f then takes the simple form $\nabla f(\boldsymbol{\alpha}) = \lambda A^T A \boldsymbol{\alpha} - \mathbf{b} = \lambda A^T \mathbf{w} - \mathbf{b}$; its (i, \mathbf{y}) -th component is $-\frac{1}{n} H_i(\mathbf{y}; \mathbf{w})$, cf. (??). Finally, note that the domain $\mathcal{D} \subset \mathbb{R}^m$ of (4) is the product of n probability simplices, $\mathcal{D} := \Delta_{|\mathcal{Y}_1|} \times \dots \times \Delta_{|\mathcal{Y}_n|}$.

TALK ABOUT DBCFW: description + pseudo-code

3. Software package

The use of the DISSOLVE^{struct} requires users to implement the following functions:

- The joint feature map $\Psi(X, Y)$ which encodes the input/output pairs.
- The structured loss function $\Delta(Y_i, Y)$.
- A maximization oracle which computes the most violating constraint by solving Eq. X.

- A prediction function that computes X . This operation is usually performed with the maximization oracle.

TODO: Can we show a simple example like the pystruct paper?

4. Experiments

Run experiments on cov and other dataset. Compare to other frameworks...

Acknowledgments

We would like to thank Jan Deriu, Bettina Messmer, Thijs Vogels and Ruben Wolff for contributing to the code and discussions to improve readability.

Appendix A.

In this appendix we prove the following theorem from Section 6.2:

Theorem *Let u, v, w be discrete variables such that v, w do not co-occur with u (i.e., $u \neq 0 \Rightarrow v = w = 0$ in a given dataset \mathcal{D}). Let N_{v0}, N_{w0} be the number of data points for which $v = 0, w = 0$ respectively, and let I_{uv}, I_{uw} be the respective empirical mutual information values based on the sample \mathcal{D} . Then*

$$N_{v0} > N_{w0} \Rightarrow I_{uv} \leq I_{uw}$$

with equality only if u is identically 0. ■

Proof. We use the notation:

$$P_v(i) = \frac{N_v^i}{N}, \quad i \neq 0; \quad P_{v0} \equiv P_v(0) = 1 - \sum_{i \neq 0} P_v(i).$$

These values represent the (empirical) probabilities of v taking value $i \neq 0$ and 0 respectively. Entropies will be denoted by H . We aim to show that $\frac{\partial I_{uv}}{\partial P_{v0}} < 0 \dots$

Remainder omitted in this sample. See <http://www.jmlr.org/papers/> for full paper.

References

- Martin Jaggi, Virginia Smith, Martin Takáč, Jonathan Terhorst, Sanjay Krishnan, Thomas Hofmann, and Michael I Jordan. Communication-Efficient Distributed Dual Coordinate Ascent. In *NIPS 2014 - Advances in Neural Information Processing Systems 27*, pages 3068–3076, 2014.
- Simon Lacoste-Julien, Martin Jaggi, Mark Schmidt, and Patrick Pletscher. Block-Coordinate Frank-Wolfe Optimization for Structural SVMs. In *ICML 2013 - Proceedings of the 30th International Conference on Machine Learning*, 2013.
- Ben Taskar, Carlos Guestrin, and Daphne Koller. Max-Margin Markov Networks. In *NIPS 2014 - Advances in Neural Information Processing Systems 27*, 2003.
- Ioannis Tsochantaridis, Thorsten Joachims, Thomas Hofmann, and Yasemin Altun. Large Margin Methods for Structured and Interdependent Output Variables. *The Journal of Machine Learning Research*, 6, December 2005.